



Easy RAG with LangChain4J

Julien Dubois
@juliendubois
Principal Manager, Java Developer Relations



What is LangChain4J?

- A Java version of LangChain, a JavaScript + Python framework
- Aims to simplify the creation of applications using Large Language Models
- Fully Open Source



Why would you use LangChain4J?

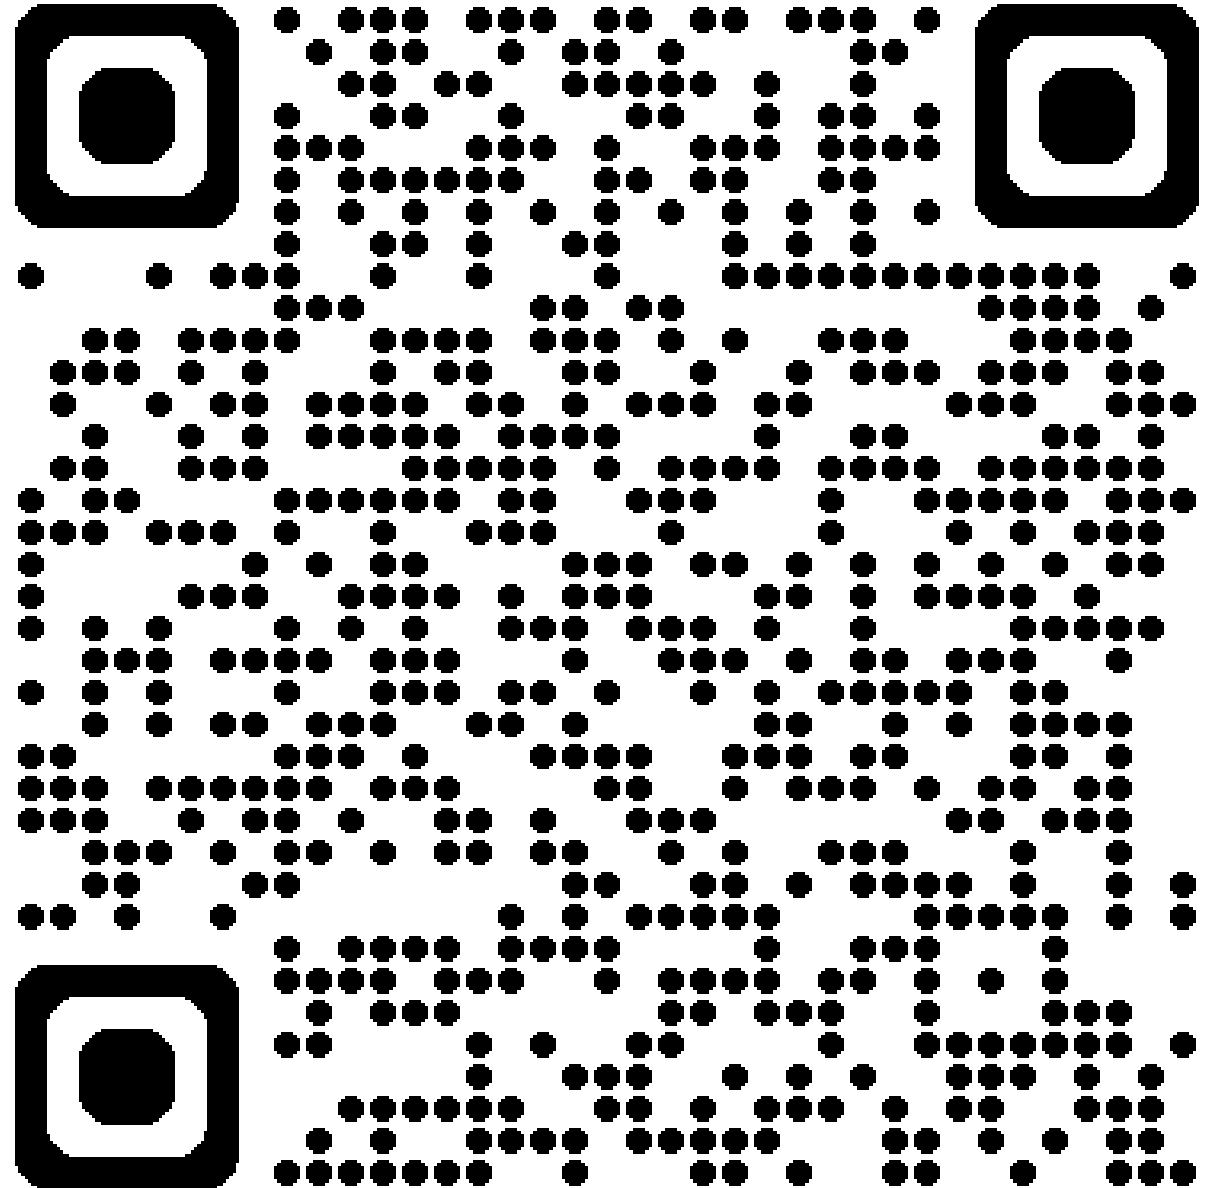
- A layer of abstraction above many APIs
- Easy integration of those APIs to build more complex applications
- Nice tooling to get work done faster



Demo code

The application used in this demo is available here:

<https://github.com/jdubois/jdubois-langchain4j-demo>



RAG = Retrieval-Augmented Generation

- 2 main limits with LLMs
 - They have a cut-off date
 - They don't have your company internal data
- How do we typically provide this data to a LLM?
 - Get new/private data and transform them into vectors
 - Store those into a vector database
 - Request the vector database and send that data to the LLM
- This is where LangChain4J excels
 - It allows to easily use vector databases and LLMs
 - It gives you access to advanced features like Hybrid Search and Semantic re-ranking (with Azure AI Search)

Let's do this!



Ingestion (1/3) – getting the data

- Document Loader
 - Fetch the data from a URL, Blob storage, GitHub repository...
- Document Transformer
 - Transforms the document to get only relevant data
- Document Splitter
 - Splits the document into segments
 - Segments can be paragraphs, lines, sentences, words... Or "recursive" to get the best one
 - Maximum segment size
 - Maximum overlap size
 - <https://langchain-text-splitter.streamlit.app>

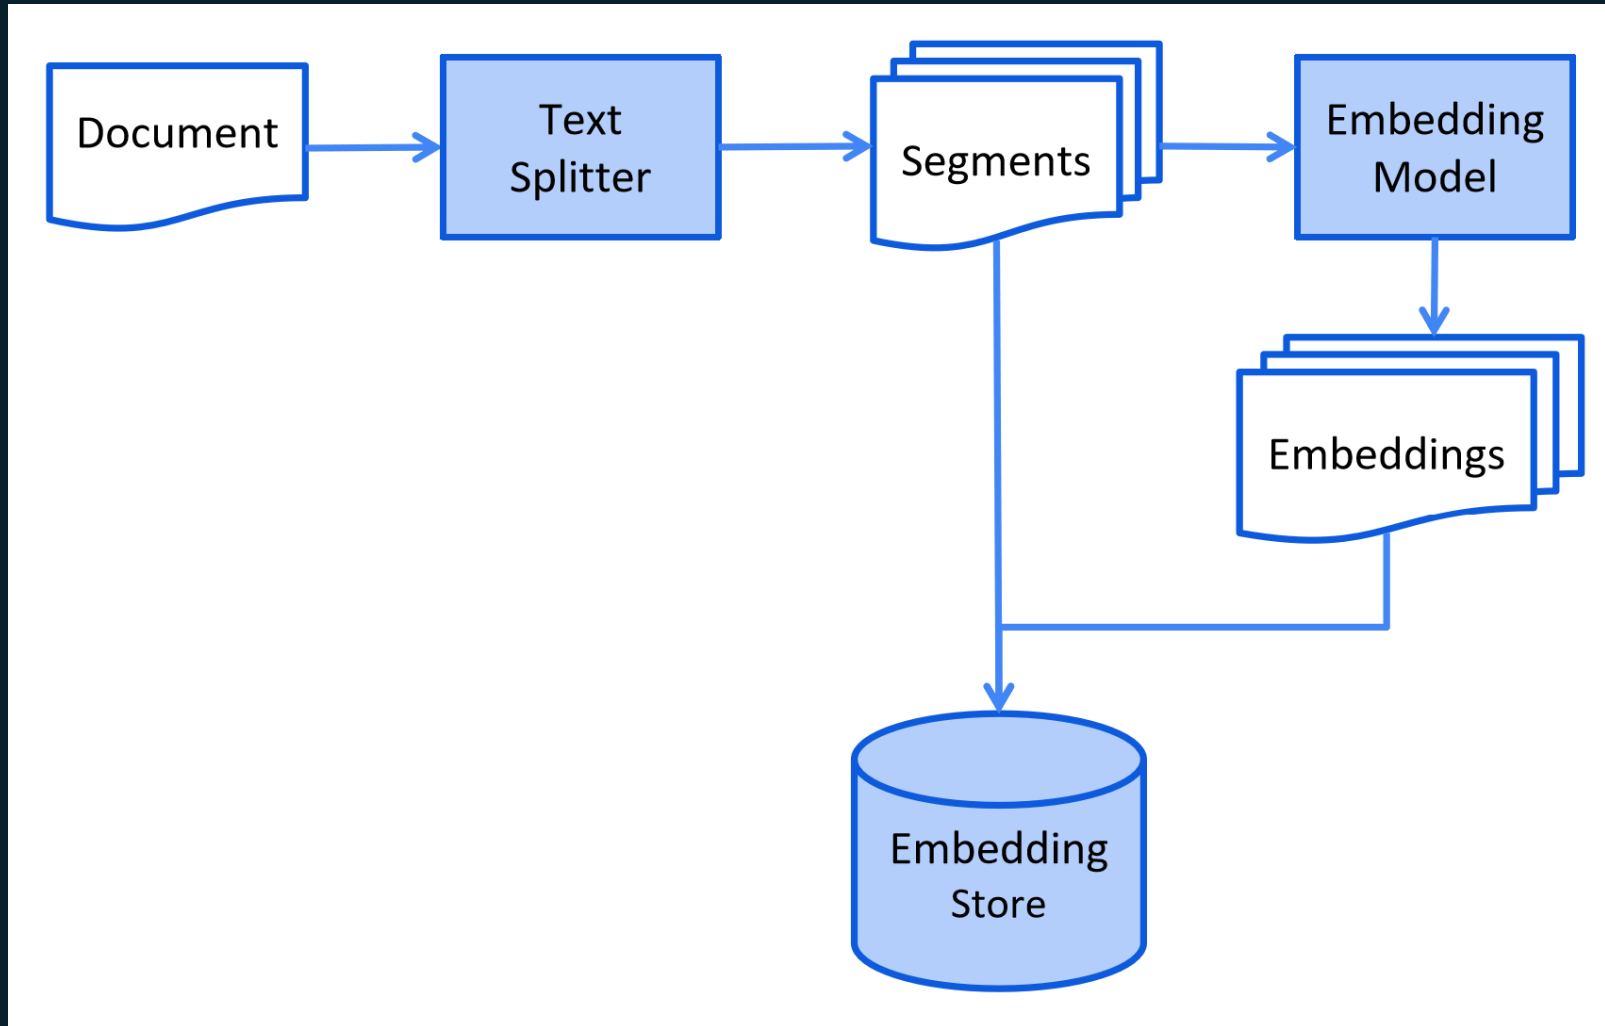
Ingestion (2/3) – transforming the text into a vector

- Embedding Model
 - Transforms the segment into a vector: an array of floats in the form [-0.021544598,-0.0643939,-0.08963279,-0.002898384,0.07627905,-0.02540857,-0.031227088,-0.0407423,-0.000047008543,-0.02396834]
 - This vector has a dimension: the size of the array (the biggest = the better & the slowest)
- There are many embedding models: "text-embedding-ada", "text-embedding-3-small", "nomic-embed-text"

Ingestion (3/3) – storing the vectors

- Embedding Store
 - Usually a vector store
 - More complex/advanced stores exist, for example adding full text search (=Lucene), and being able to mix those search mechanisms for finding better segments
- Stores all the vectors, as well as metadata
 - Metadata is very important: it should include the segment text, and usually references like file names (so sources can be quoted in the retrieval phase)
- Can find quickly vectors that point in the same direction, usually using an algorithm called “cosine similarity”

Graphical representation



Putting it all together with Easy RAG

```
Document document =  
    UrlDocumentLoader.load("https://www.microsoft.com/investor/reports/ar23/index.html",  
        new TextDocumentParser());  
  
EmbeddingStoreIngestor ingestor = EmbeddingStoreIngestor.builder()  
    .documentTransformer(new HtmlToTextDocumentTransformer(".annual-report"))  
    .documentSplitter(DocumentSplitters.recursive(300, 30))  
    .embeddingModel(embeddingModel)  
    .embeddingStore(embeddingStore)  
    .build();  
  
ingestor.ingest(document);
```

Best practices

- Clean up the text
- Test with different splitters, segment size, overlap size
- Test with different embedding models and different dimensions
- Use metadata, at least for quoting sources
- Be ready to re-index everything!
 - Have a smaller set of data for testing
 - Keep all the original data available
- Plan to index new data regularly

Retrieval-augmented generation (1/2)

- Get the user's question
- Transform it into an embedding
 - This should use the same model and dimensions as when we stored the documents
- Query the vector store to find which vectors are similar to the user's question
 - There are more advanced solutions, to help find better vectors
 - Getting too many vectors will lead to worse generation: the goal is to have the best vectors, not many vectors

Retrieval-augmented generation (2/2)

- Augment the user's question using the text segments
- Query the LLM with the augmented prompt
 - Many Chat Models are available
 - This calculation benefits more from having a GPU
- Some interesting LLMs to test with
 - GPT-4o and GPT-4o-mini : latest and best models from OpenAI
 - Phi-3.5 : small and smart model from Microsoft
 - Llama 3.2: very small model from Meta (warning: cannot be used in Europe)
 - Tinyllama: incredibly small model compatible with Llama, good for testing on a laptop

Default prompt used by EasyRAG

```
{{user question}}
```

Answer using the following information:

```
{{text segment 1}}
```


```
{{text segment 2}}
```

```
...
```

```
{{text segment n}}
```

- Possible improvements
 - Given the following information and not prior knowledge, answer the question above.
 - Keep your answer grounded in the facts of the information below.
 - If the information below doesn't contain the facts to answer, answer that you don't know.

Graphical representation



Putting it all together with Easy RAG

```
String question = "How many people are employed by Microsoft in the US?";
```

```
Assistant assistant = AIServices.builder(Assistant.class)  
    .chatLanguageModel(chatLanguageModel)  
    .contentRetriever(  
        new EmbeddingStoreContentRetriever(embeddingStore, embeddingModel, 3))  
    .build();
```

```
String answer = assistant.chat(question);
```

Best practices

- Do not send too much data in the prompt
 - This can lead to very bad answers
 - It costs tokens (=money)
- Test the Chat Models
 - Depending on the task, some models are better than others
 - Bigger models cost a lot more, so use them only if they are useful
- Improve the prompt
 - The prompt's efficiency depends on the model
- Add integration tests, for example with TestContainers
- Use GPUs in production


Putting it all together

- EasyRAG makes it easy to use the RAG pattern
 - Sensible defaults
 - Tooling
 - Easy to extend
- You'll have your RAG-based application available in no time
 - Then you'll need to improve the ingestion and test various LLMs
 - This will cost you time & money, as AI isn't replacing your hard work (yet)

Do you want the full experience?

This does not use EasyRAG on purpose: it's a complete 3-hour workshop to build everything the hard way.

<https://aka.ms/ws/openai-rag-quarkus>



Microsoft  developers